

Akash Mahajan

Site: akashmjn.me; Redwood City, CA

[@akashmjn](#) [LinkedIn/akashmjn](#)

EXPERIENCE

Member of Technical Staff / Tech Lead | [Contextual AI](#)

2024-2025

- System development (0→1)
 - Designed core [multimodal document understanding](#) (parsing + ETL index) used by retrieval agents
 - Ingesting O(billion) tokens of multimodal documents with complex layouts, tables, graphics, metadata
 - Owned system quality. Co-led with platform lead: architected modules to fit multi-stage async pipeline spanning CPU, GPU, and VLM API stages.
- Applied research: harness and eval design, token-efficient retrieval
 - Combining parsing tools, custom segmentation models and VLMs for auditable OCR with bbox provenance and minimal hallucination
 - Token-efficient synthesis beyond point queries via ingest-time ETL (à la [llm-wiki](#))
[\[Blog: Agentic alternative to GraphRAG\]](#) [\[Demo: Chat with 250 page PDF in Cursor\]](#)
- Tech Lead Manager
 - DRI cross-company; Interviewed candidates, Mentored team of 3
 - Critical in landing company's first multi-million \$ enterprise [contract with Qualcomm](#)

Senior Applied Scientist | [Microsoft Azure Speech](#)

2018-2023

- Model development: state-of-art transcription designed for scale [O(10 million) hrs/mo]
 - Shipped both batch and streaming models to [Azure Batch](#), [Microsoft Word](#), and [Microsoft Teams](#)
 - In particular: Optimized Conformer batch model ([Whisper](#)-comparable) at 50x realtime
- Applied research: diarized multi-speaker multi-mic transcription
 - Shipped diarized in-conference room transcription device covered by [The Verge](#)
 - Lead contributor: ASR training recipes, evaluation metrics, cross-system error analysis
- Research engineering: data pipelines, optimizing distributed training and inference
 - Speeding up O(1e20) FLOP training on low-cost V100 GPUs
 - Leveraged NVIDIA/ONNX profiling tools to fix bottlenecks in inference throughput

PROJECTS

Self-study: efficient audio generation, multi-stream dialog models | [\[github\]](#) (ongoing)

2026

- Speculative decoding with relaxed acceptance on single-codebook TTS model (Chatterbox-TTS)
- Meanflow distillation for single-step FSQ token flow decoder (Voxtral-TTS)
- Hybrid cascaded/duplex voice agent architectures (e.g. MoshiRAG)

tinydiarize: Lightweight extension of Whisper for diarization | [\[github\]](#)

2023

- Built extension of OpenAI's Whisper ASR model for speaker diarization with special tokens
- Released with [integration in whisper.cpp](#) (50k+ stars) — runnable on MacBooks/iPhones

EDUCATION

Stanford University, M.S. Management Science & Engineering

2016-2018

- Course Assistant for Machine Learning (CS229) & Deep Learning (CS230)
- Built low-cost conv-seq2seq TTS trainable with \$75 of compute; awarded [best poster](#) in [CS224n](#)

Indian Institute of Technology (IIT), Madras, B.Tech.

2011-2015

Chemical Engineering, minor: Control Systems (signal processing, stats, linear algebra)

MISC

Patents

- [US11044287B1](#): Network resilient real-time voice communication leveraging on-device speech models (granted during pandemic) Microsoft, 2021
- [WO2018020475A1](#): Electric vehicle drivetrain predictive health monitoring Ather Energy, 2018

Tech Stacks

- Contextual AI: Python/pydantic, HF Transformers, Gradio, vLLM, asyncio/Temporal
- Microsoft: Pytorch distributed, Azure ML/blob pipelines, ONNX/C++

Other

- Peer reviewer for ICASSP conference, TMLR journal, and NeurIPS, CVPR workshops 2025/26
 - Interviewed in [Mongabay News](#) for [OrcaHello](#): AI-assisted 24x7 hydrophone monitoring 2020/26
-